

Joint Training of Deep Boltzmann Machines for Classification

Ian J. Goodfellow

Aaron Courville
Université de Montréal

Yoshua Bengio

Abstract

We introduce a new method for training deep Boltzmann machines jointly. Prior methods require an initial learning pass that trains the deep Boltzmann machine greedily, one layer at a time, or do not perform well on classification tasks.

1 Deep Boltzmann machines

A deep Boltzmann machine (Salakhutdinov and Hinton, 2009) is a probabilistic model consisting of many layers of random variables, most of which are latent. Typically, a DBM contains a set of D input features v that are called the *visible units* because they are always observed during both training and evaluation. The DBM is usually applied to classification problems and thus often represents the class label with a one-of- k code in the form of a discrete-valued label unit y . y is observed (on examples for which it is available) during training. The DBM also contains several hidden units, which are usually organized into L layers $h^{(i)}$ of size $N_i, i = 1, \dots, L$, with each unit in a layer conditionally independent of the other units in the layer given the neighboring layers. These conditional independence properties allow fast Gibbs sampling because an entire layer of units can be sampled at a time. Likewise, mean field inference with fixed point equations is fast because each fixed point equation gives a solution to an entire layer of variational parameters.

A DBM defines a probability distribution by exponentiating and normalizing an energy function

$$P(v, h, y) = \frac{1}{Z} \exp(-E(v, h, y))$$

where

$$Z = \sum_{v', h', y'} \exp(-E(v', h', y')).$$

Z , the partition function, is intractable, due to the summation over all possible states. Maximum likelihood learning requires computing the gradient of $\log Z$. Fortunately, the gradient can be estimated using an MCMC procedure (Younes, 1999; Tieleman, 2008). Block Gibbs sampling of the layers makes this procedure efficient.

The structure of the interactions in h determines whether further approximations are necessary. In the pathological case where every element of h is conditionally independent of the others given the visible units, the DBM is simply an RBM and $\log Z$ is the only intractable term of the log likelihood. In the general case, interactions between different elements of h render the posterior $P(h | v, y)$ intractable. Salakhutdinov and Hinton (2009) overcome this by maximizing the lower bound on the log likelihood given by the mean field approximation to the posterior rather than maximizing the log likelihood itself. Again, block mean field inference over the layers makes this procedure efficient.

An interesting property of the DBM is that the training procedure thus involves *feedback connections* between the layers. Consider the simple DBM consisting of all binary valued units, with the energy function

$$E(v, h) = -v^T W^{(1)} h^{(1)} - h^{(1)T} W^{(2)} h^{(2)}.$$

Approximate inference in this model involves repeatedly applying two fixed-point update equations to solve for the mean field approximation to the posterior. Essentially it involves running a recurrent net in order to obtain approximate expectations of the latent variables.

Beyond their theoretical appeal as a deep model that admits simultaneous training of all components using a generative cost, DBMs have achieved excellent performance in practice. When they were first introduced, DBMs set the state of the art on the permutation-invariant version of the MNIST handwritten digit

recognition task at 0.95. (By permutation-invariant, we mean that permuting all of the input pixels prior to learning the network should not cause a change in performance, so using synthetic image distortions or convolution to engineer knowledge about the structure of the images into the system is not allowed). Recently, new techniques were used in conjunction with DBM pretraining to set a new state of the art of 0.79 % test error (Hinton *et al.*, 2012).

2 The joint training problem

Unfortunately, it is not possible to train a deep Boltzmann machine using only the variational bound and approximate gradient described above. Salakhutdinov and Hinton (2009) found that instead it must be trained one layer at a time, where each layer is trained as an RBM. The RBMs can then be modified slightly, assembled into a DBM, and the DBM may be trained with the learning rule described above.

In this paper, we propose a method that enables the deep Boltzmann machine to be jointly trained.

2.1 Motivation

As a greedy optimization procedure, layerwise training may be suboptimal. Recent small-scale experimental work has demonstrated this to be the case for deep belief networks (Arnold and Ollivier, 2012).

In general, for layerwise training to be optimal, the training procedure for each layer must take into account the influence that the deeper layers will provide. The standard training procedure simply does not attempt to be optimal, while the procedure advocated by (Arnold and Ollivier, 2012) makes an optimistic assumption that the deeper layers will be able to implement the best possible prior on the current layer’s hidden units. This approach does not work for deep Boltzmann machines because the interactions between deep and shallow units are symmetrical. Moreover, model architectures incorporating design features such as sparse connections, pooling, or factored multilinear interactions make it difficult to predict how best to structure one layer’s hidden units in order for the next layer to make good use of them.

Montavon and Müller (2012) showed that reparameterizing the DBM to improve the condition number of the Hessian results in successful generative training without a greedy layerwise pretraining step. However, this method has never been shown to have good classification performance, possibly because the reparameterization makes the features never be zero from the point of view of the final classifier.

2.2 Obstacles

Many obstacles make DBM training difficult. As shown by Montavon and Müller (2012), the condition number of the Hessian is poor when the model is parameterized as having binary states.

Many other obstacles exist. The intractable objective function and the great expense of methods of approximating it such as AIS makes it too costly to do line searches or early stopping. The standard means of approximating the gradient are based on stateful MCMC sampling, so any optimization method that takes large steps makes the Markov chain and thus the subsequent gradient estimates invalid.

3 The JDBM criterion

Our basic approach is to use a deterministic criterion so that each of the above obstacles ceases to be a problem.

Our specific deterministic criterion we call the *Joint DBM inpainting* criterion, given by

$$J(v, \theta) = \sum_i \log Q_i^*(v_{S_i})$$

where

$$Q^*(S_i) = \operatorname{argmin}_Q D_{KL}(Q(v_{S_i}) \| P(h | v_{-S_i})).$$

This can be viewed as a mean field approximation to the generalized pseudolikelihood. We backprop through the minimization of Q , so this can be viewed as training a family of recurrent nets that all share parameters but each optimize a different task.

While both pseudolikelihood and likelihood are asymptotically consistent estimators, their behavior in the limited data case is different. Maximum likelihood should be better for drawing samples, but generalized pseudolikelihood can often be better for training a model to answer queries conditioning on sets similar to the S_i used during training. We view our work as similar to (Stoyanov *et al.*, 2011). The idea is to train the DBM to be a general question answering machine, using the same approximations at train time as will be required at test time, rather than to train it to be a good at generating MCMC samples that resemble the training data.

We train using nonlinear conjugate gradient descent on large minibatches of data. For each data point, in the minibatch, we sample only one subset S_i to train on, rather than attempting to sum over all subsets S_i . We choose each variable in the model to be con-

ditioned on independently from the others with probability p . High values of p work best, since the mean field assumption is applied to the variables that are not selected to be conditioned on, and the more of those there are the worse the mean field assumption is.

3.1 MNIST experiments

We used the MNIST dataset as a benchmark to compare our training method to the layerwise method proposed by Salakhutdinov and Hinton (2009). In order to replicate their technique as closely as possible we refer to the accompanying demo code (<http://www.mit.edu/~rsalakhu/DBM.html>) rather than the paper itself. Since many important details of the code are not included in the paper, we provide a summary of the code here.

3.1.1 Prior method

The demo code trains a DBM consisting of v , $h^{(1)}$, $h^{(2)}$, and y . This is accomplished in three steps: 1) Training an RBM consisting of v and $h^{(1)}$ to maximize the likelihood of v . 2) Training an RBM consisting of $h^{(1)}$, $h^{(2)}$, and y to maximize the likelihood of y and $h^{(1)}$ when $h^{(1)}$ is drawn from the first RBM's posterior. 3) Assembling the RBMs into a DBM and training it to maximize the variational lower bound on $\log P(v, y)$.

Thus far the model has only been trained generatively, though the labels y are included. Its discriminative performance—its ability to predict y from v is thus somewhat limited. We used mean field inference to approximate $P(y | v)$ in the trained model and obtained a test set error of 2.15 %.

In order to obtain better discriminative performance, the DBM is used to define a feature extractor / classifier pipeline.

First, the dataset is augmented with features ϕ . ϕ is computed once at the start of discriminative training and then fixed, i.e., the discriminative learning does not change the value of ϕ . $\phi(v)$ is defined to be the mean field parameter vector $\hat{h}^{(2)}$ obtained by running mean field on v with \hat{y} clamped to 0. No explanation is given for clamping \hat{y} to 0 in the code or the paper, but we observe that it greatly improves generalization performance, even though it does not correspond to a standard probabilistic operation like marginalizing out y .

Next, these features are fed into a multilayer percep-

tron that resembles one more step of inference:

$$\begin{aligned}\hat{h}^{(1)'} &= \sigma \left(v^T A + f^T B + b^{(1)} \right) \\ \hat{h}^{(2)'} &= \sigma \left(\hat{h}^{(1)'}{}^T C + b^{(2)} \right) \\ \hat{y} &= \text{softmax} \left(\hat{h}^{(2)'}{}^T D \right)\end{aligned}$$

A , B , C , and D are initialized to $W^{(1)}$, $W^{(2)T}$, $W^{(2)}$, and $W^{(3)}$, respectively. They are then treated as independent parameters, i.e., C is not constrained to remain equal to the transpose of D during learning. The MLP is finally trained to maximize the log probability of y under \hat{y} using 100 epochs of nonlinear conjugate gradient descent.

3.2 Our method

We follow the pre-existing procedure as closely as possible. The differences are as follows:

1. We do not have a layerwise pretraining phase.
2. When training the DBM over v , $h^{(1)}$, $h^{(2)}$ and y , we use the JDBM inpainting criterion instead of PCP.
3. Rather than running training for a hard-coded number of epochs as in the DBM demo, we use early stopping based on the validation set error. We use the first 50,000 training examples for training and the last 10,000 for validation. After the validation set error starts to increase, we train on the entire MNIST training set until the log likelihood on the last 10,000 examples matches the log likelihood on the first 50,000 at the time that the validation set error began to rise.

We obtain a test set accuracy of 1.19 % on MNIST.

We observe that a DBM trained with layerwise RBM pretraining followed by standard DBM variational learning obtains a lower inpainting error on the training set than our models jointly trained using the inpainting criterion. This suggests that our criterion correctly ranks models according to their value as a classifier, but that our optimization procedure needs to be improved.

For comparison, our best result using standard DBM variational learning but without layerwise pretraining was 1.69 % test error. Using the centering trick, this increased to 2.03 %. Both of these numbers are likely to improve somewhat with more hyperparameter exploration.

References

- Arnold, L. and Ollivier, Y. (2012). Layer-wise learning of deep generative models. *ArXiv e-prints*.
- Hinton, G. E., Srivastava, N., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2012). Improving neural networks by preventing co-adaptation of feature detectors. *CoRR*, **abs/1207.0580**.
- Montavon, G. and Müller, K.-R. (2012). Learning feature hierarchies with centred deep Boltzmann machines. *CoRR*, **abs/1203.4416**.
- Salakhutdinov, R. and Hinton, G. (2009). Deep Boltzmann machines. In *Proceedings of the Twelfth International Conference on Artificial Intelligence and Statistics (AISTATS 2009)*, volume 8.
- Stoyanov, V., Ropson, A., and Eisner, J. (2011). Empirical risk minimization of graphical model parameters given approximate inference, decoding, and model structure. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 15 of *JMLR Workshop and Conference Proceedings*, pages 725–733, Fort Lauderdale. Supplementary material (4 pages) also available.
- Tieleman, T. (2008). Training restricted Boltzmann machines using approximations to the likelihood gradient. In W. W. Cohen, A. McCallum, and S. T. Roweis, editors, *ICML 2008*, pages 1064–1071. ACM.
- Younes, L. (1999). On the convergence of markovian stochastic algorithms with rapidly decreasing ergodicity rates. *Stochastics and Stochastic Reports*, **65**(3), 177–228.